
CSCI 4061

Introduction to Operating Systems

Instructor: Abhishek Chandra

Outline

- Socket Overview
- Socket Types
- Socket Operations
- TCP Client-Server Sockets

2

What are Sockets?

- Networking API provided by the OS
 - Enable applications to “hook” onto the network
 - Support one end of a network connection
- Special files
 - Program is given socket file descriptors
 - Read, write data to them
 - Special operations needed

3

Socket Types

- Protocol families
 - E.g.: TCP/IP, UUCP, Appletalk
- TCP/IP Family:
 - Connection-oriented: TCP Sockets
 - Connectionless: UDP Sockets

4

Socket Addresses

- What addresses do you need for communicating between two processes on two hosts?
 - Local process: local-IP addr, local-port
 - Remote process: remote-IP addr, remote-port

5

Socket Address: struct sockaddr

- Generic type: struct sockaddr
 - Cast to/from the actual structure type
- TCP/IP: struct sockaddr_in
 - Family: AF_INET/PF_INET
 - Port: 16-bit TCP/UDP port number
 - Address: 32-bit IP address structure

```
struct in_addr {
    in_addr_t s_addr;
};
```

6

Address Conversion

- inet_addr: Dotted to binary format
 - E.g.: 192.168.10.3 to its 32-bit equivalent
- inet_ntoa: Reverse function
- gethostbyname: Hostname to IP address
 - Uses DNS

7

Machine Byte Ordering

- Different machines use different byte orders
 - Big-endian: MSB at low addresses
 - E.g.: Sparc, PowerPC
 - Little-endian: MSB at high addresses
 - E.g.: x86
- What if sender is big-endian while receiver is little-endian?

8

Network Byte Order

- Common byte order used for all network data transmission
 - Big-endian
- Data being sent out on the network must be converted to network byte order and vice versa
 - Port numbers and IP addresses should also be converted

9

Network Byte Order Conversion

- Host-to-network byte order conversion
 - htons: 16-bit conversion
 - htonl: 32-bit conversion
- Reverse conversion
 - ntohs, ntohl
- Should we use conversion functions if we are programming on a big-endian machine?

10

Socket Operations

- Some generic operations
 - Socket creation
 - Data I/O
- Some dependent on
 - Connectionless vs. connection-oriented (TCP vs. UDP)
 - Client vs. server

11

Client-Server Model

- Client: Requests for a service
- Server: Receives requests, performs service and returns results
- E.g.: Web
 - Web browser (client) asks for a URL
 - Web server returns the corresponding file

12

Server

- Daemon that waits for requests
- When a request arrives
 - Handles the request
 - Performs some service
- Returns result to requesting client
- Goes back to waiting for more requests

13

TCP Sockets: Server Operations

- Create a socket
- Bind a local address/port number
- Wait for connections from clients
- Accept a connection
- Read request, service it, return results
- Close connection
- Service more client connections

14

Client

- Sends a request to a server
- Waits for response
- Receives response (or error)
- Done or sends more requests

15

TCP Sockets: Client Operations

- Create a socket
- Connect to a remote server
- Send request
- Receive results
- Close socket

16

Creating a Socket: `socket`

```
int socket(int family, int type, int protocol);
```

- Returns a file descriptor
 - Identifier for the socket
- Parameters
 - `family`: Protocol family. E.g.: Internet or TCP/IP (`AF_INET/PF_INET`)
 - `type`: Protocol type
 - `SOCK_STREAM`: Stream (TCP)
 - `SOCK_DGRAM`: Datagram (UDP)
 - `SOCK_RAW`: Raw (IP)
 - `protocol`: Typically 0

17

Binding a Local Address/Port: `bind`

```
int bind(int sockfd, struct sockaddr *myaddr,  
        socklen_t addrlen);
```

- Binds a local address/port based on values specified in the address structure
- Parameters:
 - `sockfd`: socket file descriptor
 - `myaddr`: address structure containing local address/port
 - `addrlen`: Length of address structure

18

Local Ports

- Well-known ports:
 - Process specifies a non-zero port number
 - Servers typically do this
 - E.g.: Web: 80, ftp: 21, ssh: 22
- Ephemeral ports:
 - Port number specified in address struct is 0
 - Kernel chooses an unused port number from a range
 - Clients typically do this

19

Local IP Addresses

- Process may specify a local IP address
 - One of the valid network interface addresses
 - Communication would happen through the chosen address
- Process may specify a wildcard IP address
 - `INADDR_ANY`
 - Kernel will choose a default IP address

20

Converting to Server Socket: listen

```
int listen(int sockfd, int backlog);
```

- Converts a socket to a "passive" server socket
 - Called only by a TCP server
- Parameters:
 - sockfd: socket file descriptor
 - backlog: Number of pending client connections

21

Accepting Connections: accept

```
int accept(int sockfd, struct sockaddr *cliaddr,  
socklen_t *addrlen);
```

- Accepts a client connection
 - Called by a TCP server after listen
- Parameters:
 - sockfd: listening (server) socket file descriptor
 - cliaddr: Client's address structure
 - addrlen: Length of address structure

22

Accepting Connections: accept

- Returns a new socket file descriptor
 - Corresponds to the TCP connection with the client
 - All communication with client happens on this new connection
- Listening socket is used only for accepting new connections
- If no new connection, server blocks on accept

23

Connecting to a Server: connect

```
int connect(int sockfd, struct sockaddr  
*servaddr, socklen_t *addrlen);
```

- Connects to a remote server at specified address and port
 - Called by a TCP client
- Parameters:
 - sockfd: socket file descriptor
 - servaddr: Server's address structure
 - addrlen: Length of address structure

24

TCP Socket I/O

- read, write
 - Used like for regular file I/O
- Remember that the return value may be different than num_bytes specified
 - May need to call read/write multiple times
- Avoid using stream operations such as fprintf, fread, etc.
 - Buffering may cause problems

25

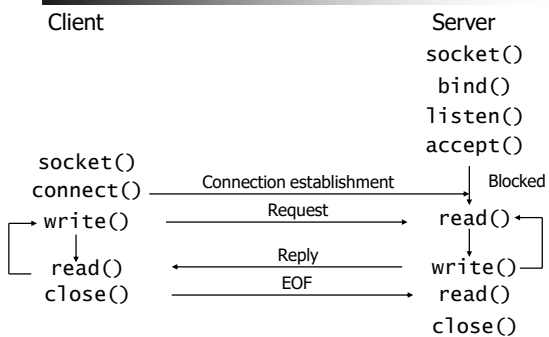
Closing a Connection: close

```
int close(int sockfd);
```

- Similar to file close
- In addition:
 - Closes TCP connection
 - Sends out any pending data before closing

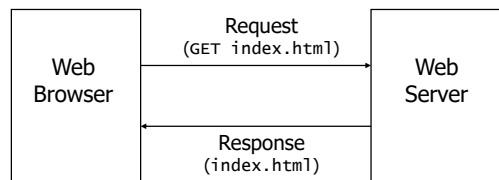
26

TCP Client-Server Operations



27

TCP Client-Server Example: Web



- Which of the client-server operations are:
 - The same?
 - Specific to Web application?

28

Handling Server Concurrency

- TCP Server has to do multiple things
 - Listen for new connections
 - Service existing client requests
 - Perform I/O on existing client connections
- Approach 1: Iterative Server
 - Do one thing at a time
 - Accept a connection, service request, close connection, go back to waiting for new connections

29

Concurrent TCP Server

- Use processes/threads for concurrency
- Main process/thread
 - Wait for new connections
 - Accept a new connection and pass on to a worker process/thread
 - Go back to waiting
- Worker processes/threads
 - Receive client connection from main process/thread
 - Service client request, perform I/O
 - Close client connection
- Can also use asynchronous I/O for concurrency

30